

# Machine Learning & Pattern Recognition

SONG Xuemeng

[songxuемeng@sdu.edu.cn](mailto:songxuемeng@sdu.edu.cn)

<http://xuемeng.bitcron.com/>

# **Generative Model and Discriminative Model**

**Imagine your task is to classify a speech to a language (Chinese, English or French).**

**Imagine your task is to classify a speech to a language (Chinese, English or French).**

1. Learning each language, and then classifying it using the knowledge you just gained.

**Imagine your task is to classify a speech to a language (Chinese, English or French).**

1. Learning each language, and then classifying it using the knowledge you just gained.

2. Determining the difference in the linguistic models without learning the languages, and then classifying the speech.

**Imagine your task is to classify a speech to a language (Chinese, English or French).**

1. Learning each language, and then classifying it using the knowledge you just gained. **Generative approach**

2. Determining the difference in the linguistic models without learning the languages, and then classifying the speech. **Discriminative approach**

# Generative vs. Discriminative

- A **generative algorithm** models how the data was generated in order to categorize a signal.
  - It asks the question: based on my generation assumptions, which category is most likely to generate this signal?
- A **discriminative algorithm** does not care about how the data was generated, it simply categorizes a given signal.

# Generative vs. Discriminative

- A **generative algorithm** models how the data was generated in order to categorize a signal.
    - It asks the question: based on my generation assumptions, which category is most likely to generate this signal?
  - A **discriminative algorithm** does not care about how the data was generated, it simply categorizes a given signal.
- 
- **Generative models** model the **distribution** of individual classes.
  - **Discriminative models** learn the **(hard or soft) boundary** between classes.



# Generative vs. Discriminative

$$P(Y|X) = \frac{P(X|Y) P(Y)}{P(X)}$$

- Generative Classifiers learn a model of the **joint** probability  $P(X, Y)$ , and make their predictions by using Bayes rules.

# Generative vs. Discriminative

$$P(Y|X) = \frac{P(X|Y) P(Y)}{P(X)}$$

- Generative Classifiers learn a model of the **joint** probability  $P(X, Y)$ , and make their predictions by using Bayes rules.
  - Assume some functional form for  $P(X|Y), P(X)$ .
  - Estimate parameters of  $P(X|Y), P(X)$  directly from training data.
  - Use Bayes rule to calculate  $P(Y|X = x_i)$ .
  - Pick the most likely label  $y$ .

# Generative vs. Discriminative

$$P(Y|X) = \frac{P(X|Y) P(Y)}{P(X)}$$

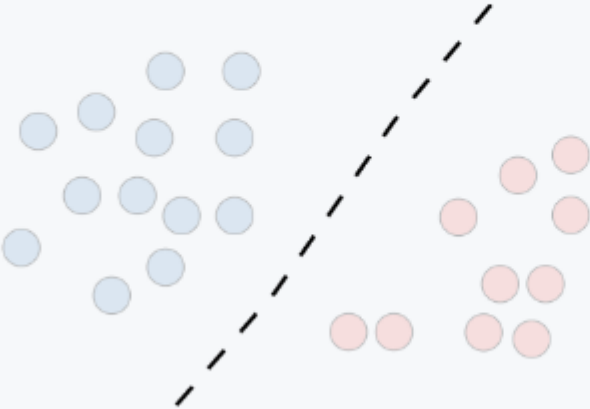
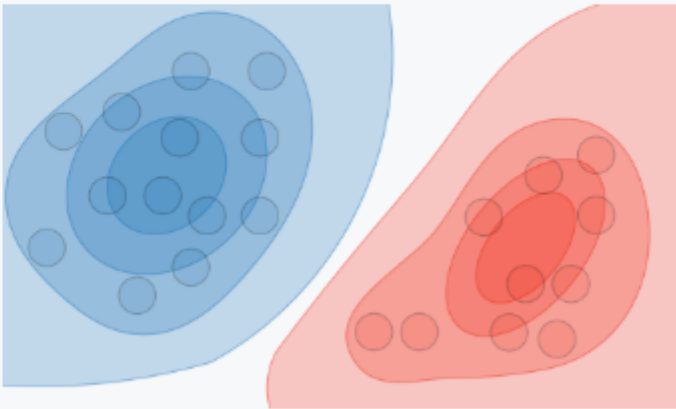
- Generative Classifiers learn a model of the **joint** probability  $P(X, Y)$ , and make their predictions by using Bayes rules.
  - Assume some functional form for  $P(X|Y), P(X)$ .
  - Estimate parameters of  $P(X|Y), P(X)$  directly from training data.
  - Use Bayes rule to calculate  $P(Y|X = x_i)$ .
  - Pick the most likely label  $y$ .
- Discriminative Classifiers model the posterior  $P(Y|X)$  directly, or learn a direct map from inputs  $x$  to the class labels  $y$ .

# Generative vs. Discriminative

$$P(Y|X) = \frac{P(X|Y) P(Y)}{P(X)}$$

- Generative Classifiers learn a model of the **joint** probability  $P(X, Y)$ , and make their predictions by using Bayes rules.
  - Assume some functional form for  $P(X|Y), P(X)$ .
  - Estimate parameters of  $P(X|Y), P(X)$  directly from training data.
  - Use Bayes rule to calculate  $P(Y|X = x_i)$ .
  - Pick the most likely label  $y$ .
- Discriminative Classifiers model the posterior  $P(Y|X)$  directly, or learn a direct map from inputs  $x$  to the class labels  $y$ .
  - Assume some function form for  $P(Y|X)$ .
  - Estimate parameters of  $P(Y|X)$  directly from training data.
  - No attempt to model underlying probability distributions.

# Generative vs. Discriminative

	Discriminative model	Generative model
Goal	Directly estimate $P(y x)$	Estimate $P(x y)$ to then deduce $P(y x)$
What's learned	Decision boundary	Probability distributions of the data
Illustration		
Examples	Regressions, SVMs	GMM, Naïve Bayes


# Generative vs. Discriminative

- Generative classifiers learn  $P(Y|X)$  indirectly and can get the wrong assumptions of the data distribution.
- Quoting Vapnik from Statistical Learning Theory:  
*“one should solve the classification problem directly and never solve a more general problem as an intermediate step, such as modeling  $p(x|y)$ ”.*


**In practice, discriminative classifiers outperform generative classifiers, if you have a lot of data.**

# Generative Models


- Gaussian Mixture Model and other mixture model



- Hidden Markov Model




- Latent Dirichlet Allocation (LDA)



# Discriminative Models

## Nearest Neighbor




Shakhnarovich, Viola, Darrell 2003

Berg, Berg, Malik 2005

...

## Neural Networks




LeCun, Bottou, Bengio, Haffner 1998

Rowley, Baluja, Kanade 1998

...

## Support Vector Machines




Guyon, Vapnik

Heisele, Serre, Poggio, 2001

...

## Boosting





# **Generative: Gaussian Mixture Model** **(GMM)**

# **Review: the Gaussian distribution**

# The Gaussian Distribution

- If random variable  $X$  is Gaussian, it has the following PDF:

$$p(x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-(x-\mu)^2/2\sigma^2}$$

- **Two parameters:** the mean  $\mu$  and the variance  $\sigma^2$  ( $\sigma$  is called the standard deviation).
- We will use “Gaussian” and “Normal” interchangeably.
- To save us some writing, we will write,

$$p(x) = \mathcal{N}(x; \mu, \sigma^2).$$

# Parameter Estimation for Gaussians: $\mu$

- Suppose we have i.i.d. observations  $X_1, \dots, X_N$  from a Gaussian distribution with unknown mean  $\mu$  and known variance  $\sigma^2$ .
- If we want to find the maximum likelihood estimate for  $\mu$ , then

# Parameter Estimation for Gaussians: $\mu$

- Suppose we have i.i.d. observations  $X_1, \dots, X_N$  from a Gaussian distribution with unknown mean  $\mu$  and known variance  $\sigma^2$ .

- If we want to find the maximum likelihood estimate for  $\mu$ , then

$$p(x_1, \dots, x_N) = \prod_{i=1}^N \mathcal{N}(x_i; \mu, \sigma^2) = \prod_{i=1}^N \frac{1}{\sqrt{2\pi}\sigma} e^{-(x_i - \mu)^2 / 2\sigma^2}$$

$$\ln p(x_1, \dots, x_N) = \sum_{i=1}^N \ln \frac{1}{\sqrt{2\pi}\sigma} - \frac{(x_i - \mu)^2}{2\sigma^2}$$

# Parameter Estimation for Gaussians: $\mu$

- Suppose we have i.i.d. observations  $X_1, \dots, X_N$  from a Gaussian distribution with unknown mean  $\mu$  and known variance  $\sigma^2$ .

- If we want to find the maximum likelihood estimate for  $\mu$ , then

$$p(x_1, \dots, x_N) = \prod_{i=1}^N \mathcal{N}(x_i; \mu, \sigma^2) = \prod_{i=1}^N \frac{1}{\sqrt{2\pi}\sigma} e^{-(x_i - \mu)^2 / 2\sigma^2}$$

$$\ln p(x_1, \dots, x_N) = \sum_{i=1}^N \ln \frac{1}{\sqrt{2\pi}\sigma} - \frac{(x_i - \mu)^2}{2\sigma^2}$$


$$\frac{d}{d\mu} \ln p(x_1, \dots, x_N) = \sum_{i=1}^N \frac{x_i - \mu}{\sigma^2} = 0 \Rightarrow \hat{\mu} = \frac{1}{N} \sum_{i=1}^N x_i$$

Does not depend on the  $\sigma^2$ .


# Gaussian Mixture Models

- GMM is useful for modeling data that from one of several groups: data points within the same group can be well-modeled by a Gaussian distribution.
- Example 1  
Suppose the **price** of a randomly chosen **paperback** book is normally distributed with mean \$10.00 and standard deviation \$1.00. Similarly, the **price** of a randomly chosen **hardback** is normally distributed with mean \$17 and variance \$1.50.

**Paperback**




**Hardback**



Is the price of randomly chosen book normally distributed?

# Gaussian Mixture Models


- GMM is useful for modeling data that from one of several groups: data points within the same group can be well-modeled by a Gaussian distribution.
- Example 2  
Suppose the **height** of a randomly chosen **man** is normally distributed with a mean around 5'9.5" and standard deviation around 2.5". Similarly, the **height** of a randomly chosen **woman** is normally distributed with a mean around 5'4.5" and standard deviation around 2.5".




Is the height of a randomly chosen person normally distributed?



# Gaussian Mixture Models




(a) Probability density for paperback books (red), hardback books (blue) and all books (black, solid)




(b) Probability density for heights of women (red), heights of men (blue) and all heights (black, solid)

Figure 1. Two Gaussian mixture models: the component densities are shown in dotted red and blue lines, while the overall density is shown as solid black line.

# Gaussian Mixture Models



(a) Probability density for paperback books (red), hardback books (blue) and all books (black, solid)



(b) Probability density for heights of women (red), heights of men (blue) and all heights (black, solid)

Figure 1. Two Gaussian mixture models: the component densities are shown in dotted red and blue lines, while the overall density is shown as solid black line.

Let's look at this a little more formally.

# Gaussian Mixture Models

- GMM is able to form smooth approximations to arbitrarily shaped densities (GMM is a universal approximator of densities).
- An arbitrary density  $f(\cdot)$ , can be approximated by a Gaussian mixture model,


$$g_k(\cdot; \boldsymbol{\omega}, \boldsymbol{\mu}, \boldsymbol{\sigma}) = \sum_{i=1}^k \omega_i \varphi(\cdot; \mu_i, \sigma_i)$$

- In the sense that

$$g_k(\cdot; \boldsymbol{\omega}, \boldsymbol{\mu}, \boldsymbol{\sigma}) \xrightarrow{k \rightarrow \infty} f(\cdot)$$


# The GMM Assumption

- There are  $k$  components. The  $i$ -th component is called  $\omega_i$
- Component  $\omega_i$  has an associated mean vector  $\mu_i$



# The GMM Assumption

- There are  $k$  components. The  $i$ -th component is called  $\omega_i$
- Component  $\omega_i$  has an associated mean vector  $\mu_i$
- Each component generates data from a Gaussian model with mean  $\mu_i$  and covariance matrix  $\sigma^2 \mathbf{I}$




# The GMM Assumption

- There are  $k$  components. The  $i$ -th component is called  $\omega_i$
- Component  $\omega_i$  has an associated mean vector  $\mu_i$
- Each component generates data from a Gaussian model with mean  $\mu_i$  and covariance matrix  $\sigma^2 \mathbf{I}$

Assume that each data point is generated as follows:

1. Pick a component at random. Choose component  $i$  with probability  $P(\omega_i)$ .




# The GMM Assumption

- There are  $k$  components. The  $i$ -th component is called  $\omega_i$
- Component  $\omega_i$  has an associated mean vector  $\mu_i$
- Each component generates data from a Gaussian model with mean  $\mu_i$  and covariance matrix  $\sigma^2 I$

Assume that each data point is generated as follows:

1. Pick a component at random. Choose component  $i$  with probability  $P(\omega_i)$ .
2. Data point  $\sim N(\mu_i, \sigma^2 I)$




# The General GMM Assumption

- There are  $k$  components. The  $i$ -th component is called  $\omega_i$
- Component  $\omega_i$  has an associated mean vector  $\mu_i$
- Each component generates data from a Gaussian model with mean  $\mu_i$  and covariance matrix  $\Sigma_i$

Assume that each data point is generated as follows:

1. Pick a component at random. Choose component  $i$  with probability  $P(\omega_i)$ .
2. Data point  $\sim N(\mu_i, \Sigma_i)$





# Computing Likelihood

Suppose we have  $x_1, x_2, \dots, x_N$  [data samples], we know priors for components  $P(\omega_1), P(\omega_2), \dots, P(\omega_k)$ , and know  $\sigma$ .

We want to find maximum likelihood estimates for  $\mu_1, \dots, \mu_k$ .

# Computing Likelihood

Suppose we have  $x_1, x_2, \dots, x_N$  [data samples], we know priors for components  $P(\omega_1), P(\omega_2), \dots, P(\omega_k)$ , and know  $\sigma$ .

We want to find maximum likelihood estimates for  $\mu_1, \dots, \mu_k$ .

This is an *unsupervised learning* problem.

We do not get to observe the label/component for each data.

# Computing Likelihood

Suppose we have  $x_1, x_2, \dots, x_N$  [data samples], we know priors for components  $P(\omega_1), P(\omega_2), \dots, P(\omega_k)$ , and know  $\sigma$ .

We want to find maximum likelihood estimates for  $\mu_1, \dots, \mu_k$ .

This is an *unsupervised learning* problem.

We do not get to observe the label/component for each data.

**Exercise**: Given the above setup, compute the log-likelihood, and then differentiate with respect to  $\mu_i$ .

# Computing Likelihood

Suppose we have  $x_1, x_2, \dots, x_N$  [data samples], we know priors for components  $P(\omega_1), P(\omega_2), \dots, P(\omega_k)$ , and know  $\sigma$ .


We want to find maximum likelihood estimates for  $\mu_1, \dots, \mu_k$ .

This is an *unsupervised learning* problem.

We do not get to observe the label/component for each data.

**Exercise**: Given the above setup, compute the log-likelihood, and then differentiate with respect to  $\mu_i$ .

**CHALLENGING!**



# Computing Likelihood

**Exercise:** Given the above setup, compute the log-likelihood, and then differentiate with respect to  $\mu_i$ .

$$\begin{aligned} P(x_1, \dots, x_N | \mu_1, \dots, \mu_k) &= \prod_{i=1}^N P(x_i | \mu_1, \dots, \mu_k) \\ &= \prod_{i=1}^N \sum_{j=1}^k P(x_i | \omega_j, \mu_1, \dots, \mu_k) P(\omega_j) = \prod_{i=1}^N \sum_{j=1}^k \mathcal{N}(x_i; \mu_j, \sigma^2) P(\omega_j) \end{aligned}$$

$P(x | \omega_j, \mu_1, \dots, \mu_k)$ : Probability that an observation from component  $\omega_j$  would have value  $x$ , given component means  $\mu_1, \dots, \mu_k$

# Computing Likelihood

**Exercise:** Given the above setup, compute the log-likelihood, and then differentiate with respect to  $\mu_i$ .

$$\begin{aligned} P(x_1, \dots, x_N | \mu_1, \dots, \mu_k) &= \prod_{i=1}^N P(x_i | \mu_1, \dots, \mu_k) \\ &= \prod_{i=1}^N \sum_{j=1}^k P(x_i | \omega_j, \mu_1, \dots, \mu_k) P(\omega_j) = \prod_{i=1}^N \sum_{j=1}^k \mathcal{N}(x_i; \mu_j, \sigma^2) P(\omega_j) \\ \ln P(x_1, \dots, x_N | \mu_1, \dots, \mu_k) &= \sum_{i=1}^N \ln \left( \sum_{j=1}^k \mathcal{N}(x_i; \mu_j, \sigma^2) P(\omega_j) \right) \end{aligned}$$

# Computing Likelihood

**Exercise:** Given the above setup, compute the log-likelihood, and then differentiate with respect to  $\mu_i$ .

$$\begin{aligned} P(x_1, \dots, x_N | \mu_1, \dots, \mu_k) &= \prod_{i=1}^N P(x_i | \mu_1, \dots, \mu_k) \\ &= \prod_{i=1}^N \sum_{j=1}^k P(x_i | \omega_j, \mu_1, \dots, \mu_k) P(\omega_j) = \prod_{i=1}^N \sum_{j=1}^k \mathcal{N}(x_i; \mu_j, \sigma^2) P(\omega_j) \\ \ln P(x_1, \dots, x_N | \mu_1, \dots, \mu_k) &= \sum_{i=1}^N \ln \left( \sum_{j=1}^k \mathcal{N}(x_i; \mu_j, \sigma^2) P(\omega_j) \right) \end{aligned}$$

Without loss of generality, we only consider the mean  $\mu_j$ . Before we dive into differentiating, we note that

$$\frac{d}{d\mu} \mathcal{N}(x; \mu, \sigma^2) = \frac{x - \mu}{\sigma^2} \cdot \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}} = \mathcal{N}(x; \mu, \sigma^2) \frac{x - \mu}{\sigma^2}$$

# Computing Likelihood

**Exercise:** Given the above setup, compute the log-likelihood, and then differentiate with respect to  $\mu_i$ .

$$\ln P(x_1, \dots, x_N | \mu_1, \dots, \mu_k) = \sum_{i=1}^N \ln \left( \sum_{j=1}^k \mathcal{N}(x_i; \mu_j, \sigma^2) P(\omega_j) \right)$$

Without loss of generality, we only consider the mean  $\mu_j$ .

$$\frac{d}{d\mu_j} \ln P(x_1, \dots, x_N | \mu_1, \dots, \mu_k) = \sum_{i=1}^N \frac{d}{d\mu_j} \ln \left( \sum_{j=1}^k \mathcal{N}(x_i; \mu_j, \sigma^2) P(\omega_j) \right)$$



# Computing Likelihood

**Exercise:** Given the above setup, compute the log-likelihood, and then differentiate with respect to  $\mu_i$ .

$$\ln P(x_1, \dots, x_N | \mu_1, \dots, \mu_k) = \sum_{i=1}^N \ln \left( \sum_{j=1}^k \mathcal{N}(x_i; \mu_j, \sigma^2) P(\omega_j) \right)$$

Without loss of generality, we only consider the mean  $\mu_j$ .

$$\begin{aligned} \frac{d}{d\mu_j} \ln P(x_1, \dots, x_N | \mu_1, \dots, \mu_k) &= \sum_{i=1}^N \frac{d}{d\mu_j} \ln \left( \sum_{j=1}^k \mathcal{N}(x_i; \mu_j, \sigma^2) P(\omega_j) \right) \\ &= \sum_{i=1}^N \frac{1}{\sum_{j=1}^k \mathcal{N}(x_i; \mu_j, \sigma^2) P(\omega_j)} \cdot P(\omega_j) \mathcal{N}(x_i; \mu_j, \sigma^2) \cdot \frac{x_i - \mu_j}{\sigma^2} \\ &= 0 \end{aligned}$$

# Computing Likelihood

**Exercise:** Given the above setup, compute the log-likelihood, and then differentiate with respect to  $\mu_i$ .

$$\ln P(x_1, \dots, x_N | \mu_1, \dots, \mu_k) = \sum_{i=1}^N \ln \left( \sum_{j=1}^k \mathcal{N}(x_i; \mu_j, \sigma^2) P(\omega_j) \right)$$

Without loss of generality, we only consider the mean  $\mu_j$ .

$$\begin{aligned} \frac{d}{d\mu_j} \ln P(x_1, \dots, x_N | \mu_1, \dots, \mu_k) &= \sum_{i=1}^N \frac{d}{d\mu_j} \ln \left( \sum_{j=1}^k \mathcal{N}(x_i; \mu_j, \sigma^2) P(\omega_j) \right) \\ &= \sum_{i=1}^N \frac{1}{\sum_{j=1}^k \mathcal{N}(x_i; \mu_j, \sigma^2) P(\omega_j)} \cdot P(\omega_j) \mathcal{N}(x_i; \mu_j, \sigma^2) \cdot \frac{x_i - \mu_j}{\sigma^2} \\ &= 0 \end{aligned}$$

There's **no way** we can solve this in closed form to get a clean maximum likelihood expression!

# The EM Algorithm

- **Expectation-maximization (EM)** is a method for finding **maximum likelihood** (or maximum a posteriori) estimate of parameter(s) in statistical model, where the model depends on **unobserved latent variables**.

# The EM Algorithm

- **Expectation-maximization (EM)** is a method for finding **maximum likelihood** (or maximum a posteriori) estimate of parameter(s) in statistical model, where the model depends on **unobserved latent variables**.

## Maximum Likelihood from Incomplete Data via the *EM* Algorithm

By A. P. DEMPSTER, N. M. LAIRD and D. B. RUBIN

*Harvard University and Educational Testing Service*

[Read before the ROYAL STATISTICAL SOCIETY at a meeting organized by the RESEARCH SECTION on Wednesday, December 8th, 1976, Professor S. D. SILVEY in the Chair]

TITLE

CITED BY

YEAR

Maximum likelihood from incomplete data via the EM algorithm

55591

1977

AP Dempster, NM Laird, DB Rubin

Journal of the royal statistical society. Series B (methodological), 1-38

# The EM Algorithm

- **Expectation-maximization (EM)** is a method for finding **maximum likelihood** (or maximum a posteriori) estimate of parameter(s) in statistical model, where the model depends on **unobserved latent variables**.



A. P. Dempster (1929~)  
Harvard University  
Department of Statistics



N. M. Laird (1943~)  
Harvard School of Public Health.  
Department of Biostatistics



D. B. Rubin (1943~)  
Harvard University (retired)  
Department of Statistics

# The EM Algorithm

- EM is an **iterative** method that alternates between performing an expectation (E) step and a maximization (M) step
  - **E-step**
  - **M-step**

# The EM Algorithm

- $Z$ : **Latent variables**;  $X$ : data;  $\Theta$ : model parameters
- $LL(\Theta|X, Z) = \ln P(X, Z|\Theta)$
- $LL(\Theta|X) = \ln P(X|\Theta) = \sum_Z P(X, Z|\Theta)$  *marginal likelihood*

**E-step** computes the **expectation of the log-likelihood** evaluated using the current estimated distributions for the latent variables based on the parameters inferred from previous step

$$Q(\Theta|\Theta^t) = E_{Z|X, \Theta^t} LL(\Theta|X, Z)$$

**M-step** computes **parameters** maximizing the expected log-likelihood from the E-step. These parameter-estimates are then used to determine the distribution of the latent variables in the next E-step.

$$\Theta^{t+1} = \arg \max_{\Theta} Q(\Theta|\Theta^t)$$

# Simple Example

Let events be “grades in a class”

$$w_1 = \text{Gets an A} \quad P(A) = \frac{1}{2}$$

$$w_2 = \text{Gets a B} \quad P(B) = \mu$$

$$w_3 = \text{Gets a C} \quad P(C) = 2\mu$$

$$w_4 = \text{Gets a D} \quad P(D) = \frac{1}{2} - 3\mu \quad (0 \leq \mu \leq 1/6)$$

Assume we want to estimate  $\mu$  from data. In a given class, there were

$a$  A's

$b$  B's

$c$  C's

$d$  D's

What's the maximum likelihood estimate of  $\mu$  given  $a, b, c, d$ ?



# Trivial Statistics

$$P(A) = \frac{1}{2} \quad P(B) = \mu \quad P(C) = 2\mu \quad P(D) = \frac{1}{2} - 3\mu,$$

$$P(a, b, c, d \mid \mu) = K(\frac{1}{2})^a (\mu)^b (2\mu)^c (\frac{1}{2} - 3\mu)^d,$$

$$\log P(a, b, c, d \mid \mu) = \log K + a \log \frac{1}{2} + b \log \mu + c \log 2\mu + d \log (\frac{1}{2} - 3\mu)$$

For maximizing the likelihood,

# Trivial Statistics

$$P(A) = \frac{1}{2} \quad P(B) = \mu \quad P(C) = 2\mu \quad P(D) = \frac{1}{2} - 3\mu,$$

$$P(a, b, c, d \mid \mu) = K(\frac{1}{2})^a (\mu)^b (2\mu)^c (\frac{1}{2} - 3\mu)^d,$$

$$\log P(a, b, c, d \mid \mu) = \log K + a \log \frac{1}{2} + b \log \mu + c \log 2\mu + d \log (\frac{1}{2} - 3\mu)$$

For maximizing the likelihood, set

$$\frac{\partial \log P}{\partial \mu} = \frac{b}{\mu} + \frac{2c}{2\mu} - \frac{3d}{\frac{1}{2} - 3\mu} = 0 \quad \Rightarrow \quad \mu = \frac{b + c}{6(b + c + d)}$$

# Trivial Statistics

$$P(A) = \frac{1}{2} \quad P(B) = \mu \quad P(C) = 2\mu \quad P(D) = \frac{1}{2} - 3\mu,$$

$$P(a, b, c, d \mid \mu) = K(\frac{1}{2})^a (\mu)^b (2\mu)^c (\frac{1}{2} - 3\mu)^d,$$

$$\log P(a, b, c, d \mid \mu) = \log K + a \log \frac{1}{2} + b \log \mu + c \log 2\mu + d \log (\frac{1}{2} - 3\mu)$$

For maximizing the likelihood, set

$$\frac{\partial \log P}{\partial \mu} = \frac{b}{\mu} + \frac{2c}{2\mu} - \frac{3d}{\frac{1}{2} - 3\mu} = 0 \quad \Rightarrow \quad \mu = \frac{b + c}{6(b + c + d)}$$

If class got

A	B	C	D
14	6	9	10

Then

$$\mu = \frac{1}{10}$$

# Same Problem with Latent Information

Someone tells us that

Number of High grades (A's + B's) =  $h$

Number of C's =  $c$

Number of D's =  $d$

What is the max likelihood estimate of  $\mu$  now?

## REMEMBER

$$P(A) = \frac{1}{2},$$

$$P(B) = \mu,$$

$$P(C) = 2\mu,$$

$$P(D) = \frac{1}{2} - 3\mu$$

$$\log(h, c, d | \mu, b)$$

$$= \log K + (h - b) \log \frac{1}{2} + b \log \mu + c \log 2\mu + d \log \left( \frac{1}{2} - 3\mu \right)$$

We can answer this question circularly...

# Same Problem with Latent Information

$$\log(h, c, d | \mu, b) = \log K + (h - b) \log \frac{1}{2} + b \log \mu + c \log 2\mu + d \log \left(\frac{1}{2} - 3\mu\right)$$

We can answer this question circularly:

## Expectation Step

If we know the value of  $\mu$  we could compute the expected value of  $b$ :

$$b = \frac{\mu}{\frac{1}{2} + \mu} h$$

## Maximization Step

If we know the expected values of  $b$  we could compute the maximum likelihood value of  $\mu$ :

$$\mu = \frac{b + c}{6(b + c + d)}$$

# EM for This Problem

We begin with a guess for  $\mu$

We iterate between **Expectation** Step and **Maximization** Step to improve our estimates of  $\mu$  and  $a$  and  $b$ .

Define  $\mu(t)$  the estimate of  $\mu$  on the  $t$ -th iteration

$b(t)$  the estimate of  $b$  on  $t$ -th iteration

- Initial guess  $\mu(0)$

$$\textbf{E-step} \quad b(t) = \frac{\mu(t)}{\frac{1}{2} + \mu(t)} h = E[b|\mu(t)]$$

$$\textbf{M-step} \quad \mu(t+1) = \frac{b(t) + c}{6(b(t) + c + d)}$$

**Continue iterating until converged.**

# EM for This Problem

**Good news: Converging to local optimum is assured.**  
**Bad news: “local” optimum.**

# EM Convergence

- Convergence proof based on fact that  $\text{Prob}(\text{data} | \mu)$  must increase or remain same between each iteration [NOT STUDY HERE]
- But it can never exceed 1 [OBVIOUS]
- So it must therefore converge [OBVIOUS]


In our example,  
suppose we had

$$h = 20$$

$$c = 10$$

$$d = 10$$

$$\mu(0) = 0$$



$t$	$\mu(t)$	$b(t)$
0	0	0
1	0.0833	2.857
2	0.0937	3.158
3	0.0947	3.185
4	0.0948	3.187
5	0.0948	3.187
6	0.0948	3.187



# Back to Learning of GMM

We have unlabeled data  $x_1, x_2, \dots, x_N$

We have the prior for  $k$  components  $P(\omega_1), P(\omega_2), \dots, P(\omega_k), \sigma$   
**Hidden (Latent)** variables,  $k$ -dimensional 0-1 vectors  $z_1, \dots, z_N$ ,  
indicating which component each point is sampled from.

$$\begin{aligned} P(x_1, \dots, x_N | \mu_1, \dots, \mu_k, z_1, \dots, z_N) &= \prod_{i=1}^N P(x_i | \mu_1, \dots, \mu_k, z_1, \dots, z_N) \\ &= \prod_{i=1}^N \sum_{j=1}^k z_i^j P(x_i | \mu_1, \dots, \mu_k, \omega_j) P(\omega_j) \\ &= \prod_{i=1}^N \sum_{j=1}^k z_i^j C \exp \left( -\frac{1}{2\sigma^2} (x_i - \mu_j)^2 \right) P(\omega_j) \end{aligned}$$

# EM for GMMs

Calculate the logarithm of the likelihood,

$$\log P(x_1, \dots, x_N | \mu_1, \dots, \mu_k, z_1, \dots, z_N)$$

$$= \sum_{i=1}^N \log \sum_{j=1}^K z_i^j C \exp \left( -\frac{1}{2\sigma^2} (x_i - \mu_j)^2 \right) P(\omega_j)$$

$$= \sum_{i=1}^N \sum_{j=1}^K z_i^j \log C \exp \left( -\frac{1}{2\sigma^2} (x_i - \mu_j)^2 \right) P(\omega_j)$$

Setting 
$$\frac{\partial \log P(x_1, \dots, x_N | \mu_1, \dots, \mu_k, z_1, \dots, z_N)}{\partial \mu_j} = 0$$

We have 
$$\mu_j = \frac{\sum_{i=1}^N z_i^j x_i}{\sum_{i=1}^N z_i^j}$$

# EM for GMMs

Iteration. On the  $t$ -th iteration let our estimates be

$$\lambda_t = \{\mu_1(t), \mu_2(t), \dots, \mu_K(t)\}$$

E-step: Computes the expectation of the log-likelihood

$$P(z_i^j = 1 | \lambda_t) = \frac{P(x_i | \omega_j, \lambda_t) P(\omega_j | \lambda_t)}{P(x_i | \lambda_t)} = \frac{P(x_i | \omega_j, \mu_j(t)) P(\omega_j)}{\sum_{l=1}^K P(x_i | \omega_l, \mu_l(t)) P(\omega_l)}$$

$$\begin{aligned} & E_{P(z_i^j | \lambda_t)} (\log P(x_1, \dots, x_N | \mu_1, \dots, \mu_K, z_1, \dots, z_N)) \\ &= E_{P(z_i^j | \lambda_t)} \sum_{i=1}^N \sum_{j=1}^K z_i^j \log C \exp \left( -\frac{1}{2\sigma^2} (x_i - \mu_j)^2 \right) P(\omega_j) \\ &= \sum_{i=1}^N \sum_{j=1}^K E_{P(z_i^j | \lambda_t)} z_i^j \log C \exp \left( -\frac{1}{2\sigma^2} (x_i - \mu_j)^2 \right) P(\omega_j) \end{aligned}$$

# EM for GMMs


**Why does  $E_{P(z_i^j | \lambda_t)}(\log P(x_1, \dots, x_N | \mu_1, \dots, \mu_k, z_1, \dots, z_N))$  work?**



# EM for GMMs

**[Jensen's inequality]** Let  $f$  be a **convex** function. Then


$$f(\lambda x_1 + (1 - \lambda)x_2) \leq \lambda f(x_1) + (1 - \lambda)f(x_2) \quad \text{where } \lambda \in [0,1]$$



- In the context of probability theory, if  $X$  is a random variable,  $f$  is **convex**, then, we have  $f(E[X]) \leq E[f(X)]$
- If  $f$  is strictly **convex**, then  $E[f(X)] = f(E[X])$  holds true if and only if  $X = E[X]$  with probability 1 (i.e., if  $X$  is constant).

# EM for GMMs

How about the concave function? (E.g., log function)



$$\log(E[X]) \geq E[\log(X)]$$

# EM for GMMs

## Log-likelihood

$$\log p_{x_1, \dots, x_N}(x_1, \dots, x_N; \lambda) = \sum_{i=1}^N \log p_{x_i}(x_i; \lambda)$$

Marginalizing over  $z_i^j$  and  
introducing  $P(z_i^j | \lambda) / P(z_i^j | \lambda)$

$$= \sum_{i=1}^N \log \sum_{z_i^j} P(z_i^j | \lambda) \frac{p_{x_i}(x_i, z_i^j | \lambda)}{P(z_i^j | \lambda)}$$

# EM for GMMs

## Log-likelihood

$$\log p_{x_1, \dots, x_N}(x_1, \dots, x_N; \lambda) = \sum_{i=1}^N \log p_{x_i}(x_i; \lambda)$$

Marginalizing over  $z_i^j$  and  
introducing  $P(z_i^j | \lambda) / P(z_i^j | \lambda)$

$$= \sum_{i=1}^N \log \sum_{z_i^j} P(z_i^j | \lambda) \frac{p_{x_i}(x_i, z_i^j | \lambda)}{P(z_i^j | \lambda)}$$

Rewriting as an expectation



# EM for GMMs

## Log-likelihood

$$\log p_{x_1, \dots, x_N}(x_1, \dots, x_N; \lambda) = \sum_{i=1}^N \log p_{x_i}(x_i; \lambda)$$

Marginalizing over  $z_i^j$  and  
introducing  $P(z_i^j | \lambda) / P(z_i^j | \lambda)$

$$= \sum_{i=1}^N \log \sum_{z_i^j} P(z_i^j | \lambda) \frac{p_{x_i}(x_i, z_i^j | \lambda)}{P(z_i^j | \lambda)}$$

Rewriting as an expectation

$$= \sum_{i=1}^N \log E_{P(z_i^j | \lambda)} \frac{p_{x_i, z_i^j}(x_i, z_i^j | \lambda)}{P(z_i^j | \lambda)}$$

# EM for GMMs

## Log-likelihood

$$\log p_{x_1, \dots, x_N}(x_1, \dots, x_N; \lambda) = \sum_{i=1}^N \log p_{x_i}(x_i; \lambda)$$

Marginalizing over  $z_i^j$  and  
introducing  $P(z_i^j | \lambda) / P(z_i^j | \lambda)$

$$= \sum_{i=1}^N \log \sum_{z_i^j} P(z_i^j | \lambda) \frac{p_{x_i}(x_i, z_i^j | \lambda)}{P(z_i^j | \lambda)}$$

Rewriting as an expectation

$$= \sum_{i=1}^N \log E_{P(z_i^j | \lambda)} \frac{p_{x_i, z_i^j}(x_i, z_i^j | \lambda)}{P(z_i^j | \lambda)}$$

Using Jensen's inequality

$$\geq \sum_{i=1}^N E_{P(z_i^j | \lambda)} \log \frac{p_{x_i, z_i^j}(x_i, z_i^j | \lambda)}{P(z_i^j | \lambda)}$$

# EM for GMMs

## Log-likelihood

$$\log p_{x_1, \dots, x_N}(x_1, \dots, x_N; \lambda) = \sum_{i=1}^N \log p_{x_i}(x_i; \lambda)$$

Marginalizing over  $z_i^j$  and  
introducing  $P(z_i^j | \lambda) / P(z_i^j | \lambda)$

$$= \sum_{i=1}^N \log \sum_{z_i^j} P(z_i^j | \lambda) \frac{p_{x_i}(x_i, z_i^j | \lambda)}{P(z_i^j | \lambda)}$$

Rewriting as an expectation

$$= \sum_{i=1}^N \log E_{P(z_i^j | \lambda)} \frac{p_{x_i, z_i^j}(x_i, z_i^j | \lambda)}{P(z_i^j | \lambda)}$$

Using Jensen's inequality

$$\geq \sum_{i=1}^N E_{P(z_i^j | \lambda)} \log \frac{p_{x_i, z_i^j}(x_i, z_i^j | \lambda)}{P(z_i^j | \lambda)}$$

$P(X, Y | Z) = P(X | Y, Z) P(Y | Z)$

$$= \sum_{i=1}^N E_{P(z_i^j | \lambda)} \log p_{x_i}(x_i | z_i^j, \lambda)$$

# EM for GMMs

Why does  $E_{P(z_i^j | \lambda_t)}(\log P(x_1, \dots, x_N | \mu_1, \dots, \mu_k, z_1, \dots, z_N))$  work?

It is the **lower bound on**  $\log P(x_1, \dots, x_N | \mu_1, \dots, \mu_k)$ .

# EM for GMMs

Iteration. On the  $t$ -th iteration let our estimates be

$$\lambda_t = \{\mu_1(t), \mu_2(t), \dots, \mu_K(t)\}$$

E-step: Computes the expectation of the log-likelihood

$$\begin{aligned} E_{P(z_i^j | \lambda_t)} (\log P(x_1, \dots, x_N | \mu_1, \dots, \mu_K, z_1, \dots, z_N)) \\ = \sum_{i=1}^N \sum_{j=1}^K E_{P(z_i^j | \lambda_t)} (z_i^j) \log K \exp \left( -\frac{1}{2\sigma^2} (x_i - \mu_j)^2 \right) P(\omega_j) \end{aligned}$$

M-step: Estimate  $\mu$  given the data's component membership distributions

$$\mu_j(t+1) = \frac{\sum_{i=1}^N E_{P(z_i^j | \lambda_t)} (z_i^j) x_i}{\sum_{i=1}^N E_{P(z_i^j | \lambda_t)} (z_i^j)}$$

# EM for **General** GMMs

$P_i(t)$  is shorthand for estimate of  $P(\omega_j)$  on  $t$ -th iteration

Iteration. On the  $t$ -th iteration let our estimates be

$$\lambda_t = \{\mu_1(t), \mu_2(t), \dots, \mu_K(t), \Sigma_1(t), \Sigma_2(t), \dots, \Sigma_K(t), P_1(t), P_2(t), \dots, P_K(t)\}$$

E-step: Computes the expectation of the log-likelihood

$$P(z_i^j = 1 | \lambda_t) = \frac{P(x_i | \omega_j, \lambda_t) P(\omega_j | \lambda_t)}{P(x_i | \lambda_t)} = \frac{P(x_i | \omega_j, \Sigma_j(t), \mu_j(t)) P_j(t)}{\sum_{l=1}^K P(x_i | \omega_l, \Sigma_l(t), \mu_l(t)) P_l(t)}$$

$$E_{P(z_i^j | \lambda_t)} (\log P(x_1, \dots, x_N | \lambda_t, z_1, \dots, z_N))$$

# EM for **General** GMMs

$P_i(t)$  is shorthand for estimate of  $P(\omega_j)$  on  $t$ -th iteration

Iteration. On the  $t$ -th iteration let our estimates be

$$\lambda_t = \{\mu_1(t), \mu_2(t), \dots, \mu_K(t), \Sigma_1(t), \Sigma_2(t), \dots, \Sigma_K(t), P_1(t), P_2(t), \dots, P_K(t)\}$$

E-step: Computes the expectation of the log-likelihood

$$E_{P(z_i^j | \lambda_t)}(\log P(x_1, \dots, x_N | \lambda_t, z_1, \dots, z_N))$$

M-step : Estimate **parameters** given the data's component membership distributions

$$\frac{\partial E_{P(z_i^j | \lambda_t)}(\log P(x_1, \dots, x_N | \lambda_t))}{\partial \mu_j} = 0 \quad \frac{\partial E_{P(z_i^j | \lambda_t)}(\log P(x_1, \dots, x_N | \lambda_t))}{\partial \Sigma_j} = 0$$

$$\frac{\partial E_{P(z_i^j | \lambda_t)}(\log P(x_1, \dots, x_N | \lambda_t))}{\partial P_j} = 0$$

# EM for **General** GMMs

$P_i(t)$  is shorthand for estimate of  $P(\omega_j)$  on  $t$ -th iteration

Iteration. On the  $t$ -th iteration let our estimates be

$$\lambda_t = \{\mu_1(t), \mu_2(t), \dots, \mu_K(t), \Sigma_1(t), \Sigma_2(t), \dots, \Sigma_K(t), P_1(t), P_2(t), \dots, P_K(t)\}$$

E-step: Computes the expectation of the log-likelihood

$$E_{P(z_i^j | \lambda_t)}(\log P(x_1, \dots, x_N | \lambda_t, z_1, \dots, z_N))$$


M-step : Estimate **parameters** given the data's component membership distributions

$$\mu_j(t+1) = \frac{\sum_{i=1}^N P(z_i^j | \lambda_t) x_i}{\sum_{i=1}^N P(z_i^j | \lambda_t)} \quad P_j(t+1) = \frac{\sum_{i=1}^N P(z_i^j | \lambda_t)}{N}$$


$$\Sigma_j(t+1) = \frac{\sum_{i=1}^N P(z_i^j | \lambda_t) [x_i - \mu_j(t+1)][x_i - \mu_j(t+1)]^T}{\sum_{i=1}^N P(z_i^j | \lambda_t)}$$




# Gaussian Mixture Model Example: Start




After 1st  
iteration




# After 2nd iteration




**After 3rd  
iteration**




After 4th  
iteration




After 5th  
iteration



After 6th  
iteration




After 20th  
iteration





# How to Use GMM for Classification

- Train GMM for each class , and then use Bayesian Rule for classification



Example of binary classification using two GMMs

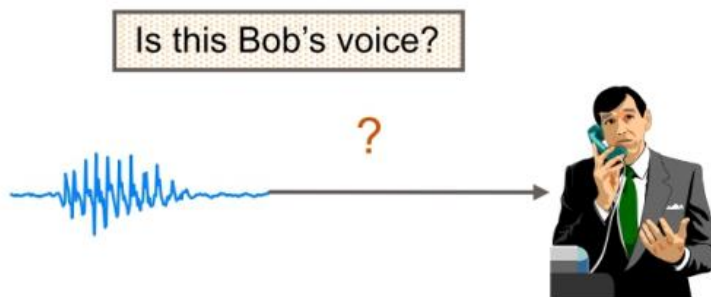
Train each GMM separately, using data set of Class 1 for GMM1 and dataset of class 2 for GMM2 (3 Gaussians for each GMM)

# How to Use GMM for Classification


- Train GMM for each class , and then use Bayesian Rule for classification
- Train universal GMM, and then adapt it for individual class, and finally do classification
  - Widely used in speaker verification

# Speaker Verification and Identification

- **Speaker Verification:** Determine whether unknown speaker matches a **specific** speaker
  - One-to-one mapping
- **Speaker Identification:** Determine whether unknown speaker matches one of a set known speakers
  - One-to-many mapping





Speaker Verification



Speaker Identification


# Speaker Verification and Identification

- **Speaker Verification:** Determine whether unknown speaker matches a **specific** speaker
  - One-to-one mapping
- **Speaker Identification:** Determine whether unknown speaker matches one of a set known speakers
  - One-to-many mapping




# GMM-UBM Speaker Verification

- A GMM, namely the **universal background model** (UBM), is trained to represent the speech of the general population.
- Speaker GMM (**Target model**) is established by adjusting UBM by using **MAP adaptation**.



# MAP Adaption


- In practice, only the mean vectors will be adapted:



# Maximum a Posteriori (MAP)

- The MAP algorithm finds the parameters of target-speaker's GMM given UBM parameters
- First step is the same as EM. Given  $T_s$  acoustic vectors  $\mathcal{X}^{(s)} = \{\mathbf{x}_1, \dots, \mathbf{x}_{T_s}\}$  from speaker  $s$ , we compute the statistics:

Probability estimated by the UBM


$$n_c = \sum_{t=1}^{T_s} \gamma_t(c) \quad \text{and} \quad E_c(\mathcal{X}^{(s)}) = \frac{1}{n_c} \sum_{t=1}^{T_s} \gamma_t(c) \mathbf{x}_t$$

- Adapt UBM parameters by

$$\mu_c^{(s)} = \alpha_c E_c(\mathcal{X}^{(s)}) + (1 - \alpha_c) \mu_c^{\text{ubm}}$$

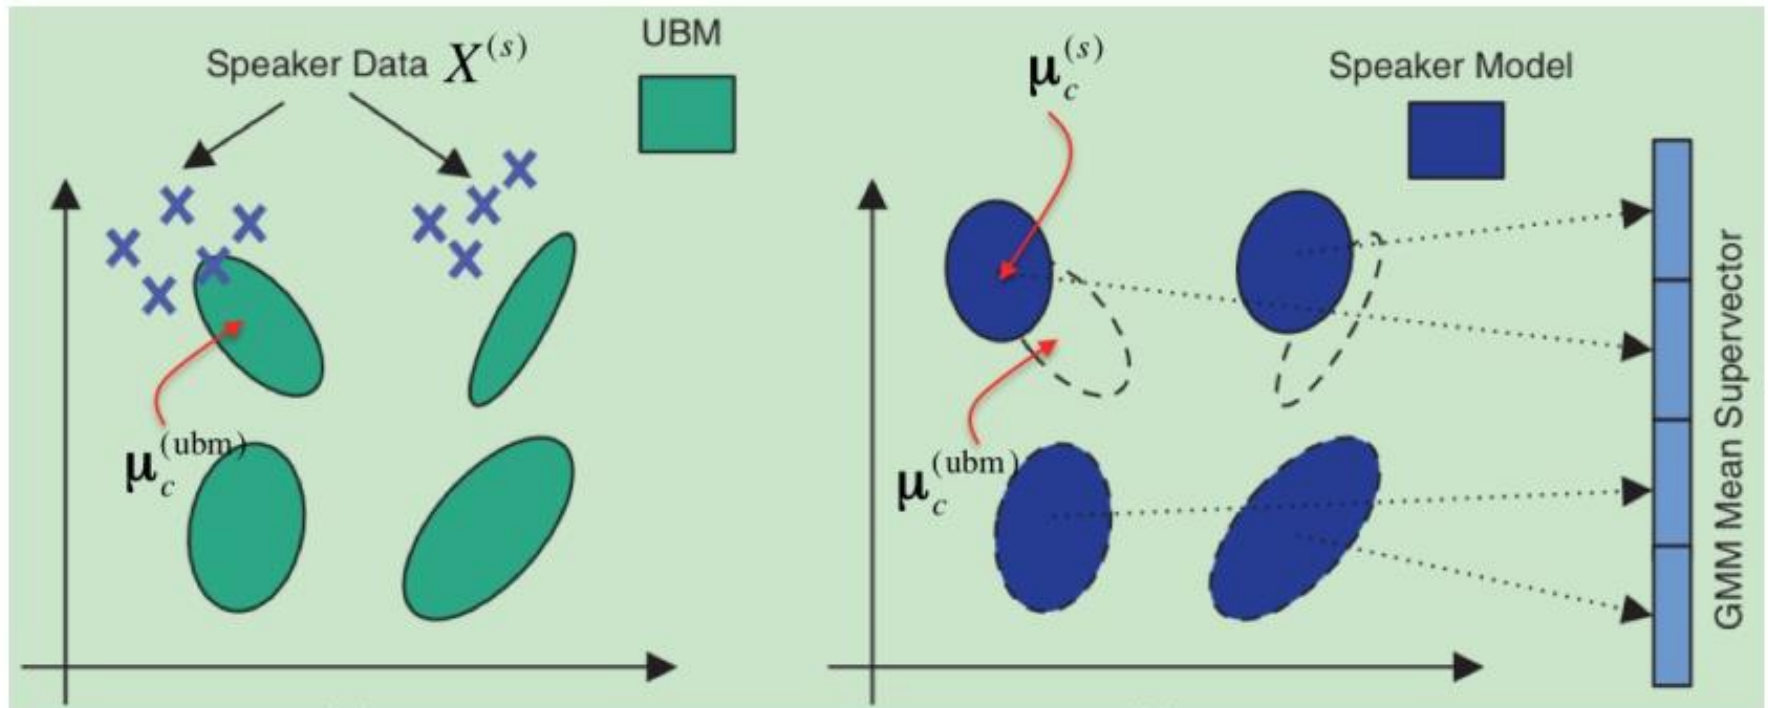
where

$$\alpha_c = \frac{n_c}{n_c + r}$$

and  $r$  is called the relevance factor. Typically,  $r = 16$ .

# MAP Adaption

- Adapt the UBM model to each speaker using the MAP algorithm:<sup>1</sup>



$$\mu_c^{(s)} = \alpha_c E_c(\mathcal{X}^{(s)}) + (1 - \alpha_c) \mu_c^{ubm}$$

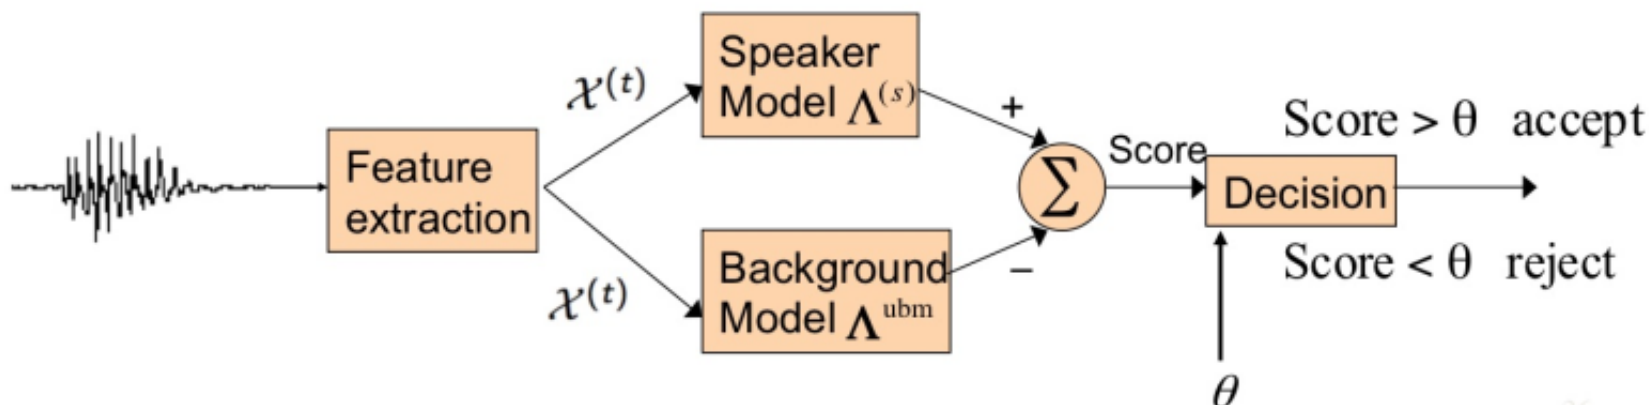
- $\alpha_c \rightarrow 1$  when  $\mathcal{X}^{(s)}$  comprises lots of data and  $\alpha_c \rightarrow 0$  otherwise.



# GMM-UBM Scoring

- Given the acoustic vectors  $\mathcal{X}^{(t)}$  from a test speaker and a claimed identity  $s$ , speaker verification can be formulated as a 2-class hypothesis problem:
  - $H_0$ :  $\mathcal{X}^{(t)}$  comes from the true speaker  $s$
  - $H_1$ :  $\mathcal{X}^{(t)}$  comes from an impostor
- Verification score is a log-likelihood ratio:

$$S_{\text{LR}}(\mathcal{X}^{(t)} | \Lambda^{(s)}, \Lambda^{\text{ubm}}) = \log p(\mathcal{X}^{(t)} | \Lambda^{(s)}) - \log p(\mathcal{X}^{(t)} | \Lambda^{\text{ubm}})$$



# Conclusion

- **Generative Models**
- **Discrimination Models**
- **GMM**
- **EM Algorithm**